

Markov Decision Processes

Riccardo Brioschi

21 February 2024

Foundations of Reinforcement Learning



Introduction

Many RL papers contain a background section like the following one:

The Option-Critic Architecture

Pierre-Luc Bacon, Jean Harb, Doina Precup
Reasoning and Learning Lab, School of Computer Science
McGill University
{pbacon, jharb, dprecup}@cs.mcgill.ca

Preliminaries and Notation

A Markov Decision Process consists of a set of states \mathcal{S} , a set of actions \mathcal{A} , a transition function $P : \mathcal{S} \times \mathcal{A} \rightarrow (\mathcal{S} \rightarrow [0, 1])$ and a reward function $r : \mathcal{S} \times \mathcal{A} \rightarrow \mathbb{R}$. For convenience, we develop our ideas assuming discrete state and action sets. However, our results extend to continuous spaces using usual measure-theoretic assumptions (some of our empirical results are in continuous tasks). A (Markovian stationary) *policy* is a probability distribution over actions conditioned on states, $\pi : \mathcal{S} \times \mathcal{A} \rightarrow [0, 1]$. In discounted problems, the value function of a policy π is defined as the expected return: $V_\pi(s) = \mathbb{E}_\pi [\sum_{t=0}^{\infty} \gamma^t r_{t+1} \mid s_0 = s]$ and its action-value function as $Q_\pi(s, a) = \mathbb{E}_\pi [\sum_{t=0}^{\infty} \gamma^t r_{t+1} \mid s_0 = s, a_0 = a]$, where $\gamma \in [0, 1)$ is the *discount factor*. A policy π is *greedy* with respect to a given action-value function Q if

In this lecture you will learn

1. What a Markov Decision Process is.
2. How MDPs can be solved with dynamic programming.
3. How future discounted MDPs can be solved with value iteration or policy iteration.

Recommended reading:
Sutton & Barto, Chapters 3 & 4

Markov Decision Processes


We define as a **Markov Decision Process (MDP)** a stochastic process characterized by:

- ▶ finite **state space** \mathcal{S} with $|\mathcal{S}| < \infty$,
- ▶ finite **action spaces** $\{\mathcal{A}_s | s \in \mathcal{S}\}$ with $|\mathcal{A}_s| < \infty$,
- ▶ **transition probabilities** $p_{s_i \rightarrow s_j}^a \in [0, 1]$ (alternatively, $P(s_j | s_i, a)$),
- ▶ **immediate rewards** $R_{s_i \rightarrow s_j}^a$ and $r_{s_i}^a = \sum_{s_j} p_{s_i \rightarrow s_j}^a \mathbb{E}[R_{s_i \rightarrow s_j}^a] \in \mathbb{R}$,
- ▶ **discount factor** $\gamma \in [0, 1]$,
- ▶ and **initial state probabilities** $p_{s_i}^{(0)}$.

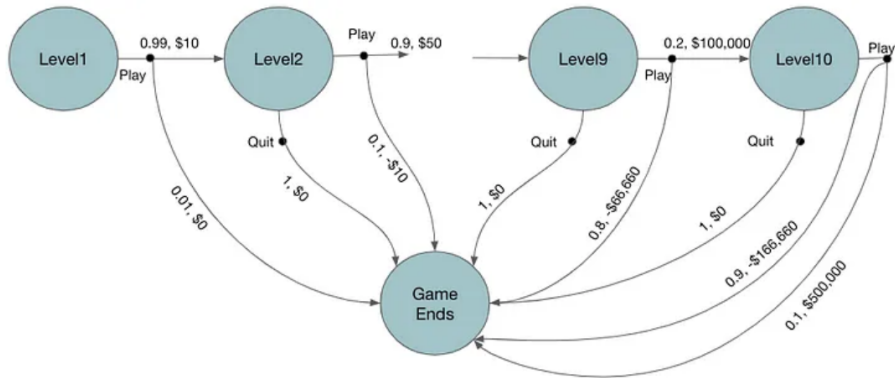
For a sequence (or trajectory) of state-action-reward tuples, we will use the notation

$$\tau = (S_0, A_0, R_1, S_1, A_1, \dots, R_T) \text{ where } R_i = R_{S_{i-1} \rightarrow S_i}^{A_{i-1}}.$$

MDPs: An Example



MDPs: An Example



Policies

A policy π selects an action based on the previous history $\tau = (S_{0:T}, A_{0:T})$

- ▶ In this course, we only consider **Markov Policies**, i.e. dependent only on the current state
- ▶ It is a mapping $\pi : \mathcal{S} \rightarrow \mathcal{A}$ or $\pi : \mathcal{S} \rightarrow \Delta(\mathcal{A})$, $\Delta(\mathcal{A})$ denotes the simplex over \mathcal{A}

Families of policies

Deterministic Policy

- ▶ Stationary policy $\pi : \mathcal{S} \rightarrow \mathcal{A}$, $a_t = \pi(s_t)$
- ▶ Markov policy $\pi_t : \mathcal{S} \rightarrow \mathcal{A}$, $a_t = \pi_t(s_t)$

Randomized Policy:

- ▶ Stationary policy $\pi : \mathcal{S} \rightarrow \Delta(\mathcal{A})$, $a_t \sim \pi(\cdot | s_t)$
- ▶ Markov policy $\pi_t : \mathcal{S} \rightarrow \Delta(\mathcal{A})$, $a_t \sim \pi_t(\cdot | s_t)$

Returns

Acting on a MDP results in immediate rewards $R_{s_i \rightarrow s_j}^a$. Accumulating these rewards, we obtain the return.

- ▶ **Finite time horizon T :** $\mathbb{E}_\pi \left[\sum_{t=1}^T \gamma^{(t-1)} R_t \right]$
- ▶ **Discounted Reward:** $J(\pi) = \mathbb{E}_\pi \left[\sum_{t=1}^{\infty} \gamma^{(t-1)} R_t \right]$

These quantities are obtained averaging the returns over the trajectories obtained moving in the MDP, according to the policy and the transition probabilities of the environment.

Q-Values and V-Values

- ▶ It is useful to introduce **Q-values** and **V-values**

$$Q_{\gamma}^{(T)}(\pi, s, a) = \mathbb{E} \left[\sum_{t=1}^T \gamma^{(t-1)} R_t \middle| S_0 = s, A_0 = a \right]$$

$$V_{\gamma}^{(T)}(\pi, s) = \mathbb{E} \left[\sum_{t=1}^T \gamma^{(t-1)} R_t \middle| S_0 = s \right]$$

- ▶ These two quantities are strongly related, indeed:

$$V_{\gamma}^{(T)}(\pi, s) = \sum_{a \in \mathcal{A}_s} \pi(a|s) Q_{\gamma}^{(T)}(\pi, s, a)$$

- ▶ These quantities can also be generalized to the infinite horizon setting

Value Functions and Objectives

The goal is to find a **policy** $\pi^{(t)}(a|s) \in [0, 1]$ that maximizes some objective. We define the horizon- T **value function (V-value)**

$$\begin{aligned} V_{\gamma}^{(T)}(\pi, s) &= \mathbb{E}_{\pi} \left[\sum_{t=1}^T \gamma^{(t-1)} R_t \middle| S_0 = s \right] \\ &= \sum_{A_0, S_1, A_1, \dots, A_{T-1}} \Pi(\tau) \left(r_s^{A_0} + \gamma r_{S_1}^{A_1} + \dots + \gamma^{T-1} r_{S_{T-1}}^{A_{T-1}} \right) \end{aligned}$$

where $\Pi(\tau) = \prod_{t=0}^{T-1} \pi^{(t)}(A_t|s_t) p_{s_t \rightarrow s_{t+1}}^{A_t}$ is the probability of trajectory τ

- ▶ Find the policy π that maximizes $V_{\gamma}^{(T)}(\pi, s)$ for all $s \in \mathcal{S}$
- ▶ **Future Discounted Values:** $V_{\gamma}^{\infty}(\pi, s) = \lim_{T \rightarrow \infty} V_{\gamma}^{(T)}(\pi, s)$ for $\gamma \in [0, 1)$.

The Optimal Fixed Horizon Policy

The policy π^* can be found with **Dynamic Programming**.


1. The optimal horizon-1 values are $V_{\gamma}^{(1)}(\pi^*, s) = \max_{a \in \mathcal{A}_s} r_s^a$.
2. The optimal horizon- $(t + 1)$ values are

$$V_{\gamma}^{(t+1)}(\pi^*, s) = \max_{a \in \mathcal{A}_s} Q_{\gamma}^{(t+1)}(\pi^*, s, a) = \max_{a \in \mathcal{A}_s} \left(r_s^a + \gamma \sum_{s' \in \mathcal{S}} p_{s \rightarrow s'}^a V_{\gamma}^{(t)}(\pi^*, s') \right)$$

This result is based on **Bellman's Principle of Optimality**.

The horizon- T policy is not stationary, in general, i.e. $\pi^{(t)}(a|s) \neq \pi^{(t')}(a|s)$ for $t \neq t'$, but it can be chosen to be deterministic. Any idea why?

Optimal Fixed Horizon Policy: An Example



The Optimal Infinite Horizon Policy

We are now looking for a policy π which maximizes

$$V_{\gamma}^{(\infty)}(\pi, s) = \mathbb{E}_{\pi} \left[\sum_{t=1}^{\infty} \gamma^{(t-1)} R_t \middle| S_0 = s \right], \forall s \in \mathcal{S}$$

- ▶ The optimal policy is now stationary
- ▶ As done before, we look for a deterministic policy, for which

$$V_{\gamma}^{(\infty)}(\pi, s) = \max_{a \in \mathcal{A}_s} \left(r_s^a + \gamma \sum_{s' \in \mathcal{S}} p_{s \rightarrow s'}^a V_{\gamma}^{(\infty)}(\pi, s') \right) \quad (1)$$


Fixed-Point Iterations and Banach's Fixed Point Theorem

Some equations of the form
 $x = T(x)$ can be solved with a
fixed point iteration:

Start with $x^{(0)}$ and compute

$$x^{(k)} = T(x^{(k-1)})$$

until $x^{(k)} \approx x^{(k-1)}$.



All we need is proving that our
expression is a contraction!

Notes

Some equations of the form
 $x = T(x)$ can be solved with a
 fixed-point iteration.
 Start with $x^{(0)}$ and compute
 $x^{(k)} = T(x^{(k-1)})$
 until $x^{(k)} \approx x^{(k-1)}$.
 All we need is proving that our
 expression is a contraction!



Let (X, d) be a complete metric space. Then a map $T : X \rightarrow X$ is called a **contraction mapping** on X if there exists $q \in [0, 1)$ such that $d(T(x), T(y)) \leq qd(x, y)$ for all $x, y \in X$.

Banach Fixed Point Theorem. Let (X, d) be a non-empty complete metric space with a contraction mapping $T : X \rightarrow X$. Then T admits a unique fixed-point x^* in X (i.e. $T(x^*) = x^*$). Furthermore, x^* can be found as follows: start with an arbitrary element $x_0 \in X$ and define a sequence $(x_n)_{n \in \mathbb{N}}$ by $x_n = T(x_{n-1})$ for $n \geq 1$. Then $\lim_{n \rightarrow \infty} x_n = x^*$.

Maximizing Future Discounted Values with Dynamic Programming

Let us define the mapping (sometimes called **Bellman operator**)

$$T_\gamma : \mathbb{R}^{|\mathcal{S}|} \rightarrow \mathbb{R}^{|\mathcal{S}|}, T_\gamma(X)_s = \max_{a \in \mathcal{A}_s} \left(r_s^a + \gamma \sum_{s' \in \mathcal{S}} p_{s \rightarrow s'}^a X_{s'} \right). \quad (2)$$

- ▶ One can show that the mapping T_γ is a contraction mapping and Banach's fixed point theorem can be applied. Hence, there is a unique fixed point $X^* = T_\gamma(X^*)$.
- ▶ The optimal policy is to choose actions in $\arg \max_{a \in \mathcal{A}_s} Q_\gamma^\infty(\pi^*, s, a)$.
- ▶ This policy is **stationary** and it can be chosen to be **deterministic**.

Value Iteration

Iteratively compute horizon- t values until $\max_{s \in \mathcal{S}} |V_{\gamma}^{(t+1)}(\pi^*, s) - V_{\gamma}^{(t)}(\pi^*, s)| < \theta$, where $\theta > 0$ is some convergence criterion. The optimal stationary policy picks actions in $\arg \max_{a \in \mathcal{A}_s} Q_{\gamma}^{t^*}(\pi^*, s, a)$, where t^* is the stopping iteration.

Value Iteration, for estimating $\pi \approx \pi_*$

Algorithm parameter: a small threshold $\theta > 0$ determining accuracy of estimation

Initialize $V(s)$, for all $s \in \mathcal{S}^+$, arbitrarily except that $V(\text{terminal}) = 0$

Loop:

```
|  $\Delta \leftarrow 0$   
| Loop for each  $s \in \mathcal{S}$ :  
|    $v \leftarrow V(s)$   
|    $V(s) \leftarrow \max_a \sum_{s', r} p(s', r | s, a) [r + \gamma V(s')]$   
|    $\Delta \leftarrow \max(\Delta, |v - V(s)|)$ 
```

until $\Delta < \theta$

Output a deterministic policy, $\pi \approx \pi_*$, such that

$$\pi(s) = \arg \max_a \sum_{s', r} p(s', r | s, a) [r + \gamma V(s')]$$

What is the relationship to Reinforcement Learning?

“Solving” an MDP amounts to solving an optimal control problem, i.e. finding the optimal policy, where the dynamics is known, i.e. $p_{s_i \rightarrow s_j}^a$ and r_s^a are assumed to be known. On the contrary, as we will see, in Reinforcement Learning, one assumes that the dynamics and rewards are unknown.

References

- ▶ <https://towardsdatascience.com/real-world-applications-of-markov-decision-process-mdp-a39685546026>
- ▶ Reinforcement Learning: An Introduction, Sutton and Barto
- ▶ Reinforcement Learning course (CS-456), EPFL